

# Optimizer

- Motivation
  - finds the optimize parameters(weights, learning rates...) that minimize the error(loss function), increasing the accuracy and training speed of the model.
- Gradient Descent(3 types)
  - Concepts Understanding
    - Learning rate: how much model weights should ne updated
    - Batch: the number of samples to be taken for updating model parameters
  - starts by quickly going down the steepest slope, then slowly goes down the bottom of the valley.
- **GSD with Momentum(GSDM)**
  - **Momentum** helps in faster convergence of the loss function
    - $m \leftarrow \beta m - \eta \delta_{\theta} L(\theta)$
    - $\theta \leftarrow \theta + m$ 
      - $m$ : momentum vector
      - $\beta$ : momentum, for avoiding the momentum going too large  $\in [0, 1]$ , normal set=0.9
  - learning rate is decreased with a high momentum term, because of the high oscillations
  - Downside:
    - momentum  $\uparrow$ , the possibility of passing the optimal minimum also increases.
    - result in poor accuracy and even more *oscillations*.
    - add another hyperparameter to tune
  - use exponential moving average of the gradients to update weights and bias, reduce the noise, and smoothen the data.
    - $w_t = w_{t-1} - \eta V_{dw_t}$ 
      - $V_{dw_t} = \beta V_{dw_{t-1}} + (1 - \beta) \frac{\delta L}{\delta w_{t-1}}$
    - $b_t = b_{t-1} - \eta V_{db_t}$ 
      - $V_{db_t} = \beta V_{db_{t-1}} + (1 - \beta) \frac{\delta L}{\delta b_{t-1}}$
      - $L$ : loss function
  - Coding: `optimizer=keras.optimizer.SGD(lr=0.001,momentum=0.9)`
    - As results,  $\beta = 0.9$ , momentum optimization 10 times faster than gradient time\* learning rate
- **Adam** (default)
  - updates the learning rate for each network parameter individually from estimates of first and second moments of the gradients.

- Adam=adagrad(works well on sparse gradients) + RMSProp(works well in online and nonstationary settings)

- reduces the radically diminishing learning rates of AdaGrad
  - Reason
    - use exponential moving average of the gradients to scale the learning rate
    - while, adagrad use simple average method; RMSProp use exponential decaying sum.

- Functions


- $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
- $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$ 
  - $m_t$ : the first moment estimation
  - $v_t$ : the second moment estimation
  - $\beta_1$  and  $\beta_2$  :  $[0, 1)$  control the exponential decay rate of moving average.
  - $g_t$ : gradient at time t along parameter  $w$
- $\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$
- $\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$ 
  - bias corrected estimates<--initialization bias for moment estimation set to be 0
- $\theta_{t+1} = \theta_t - \frac{\eta \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$

- coding: `optimizer=keras.optimizers.Adam(lr=0.001,beta_1=0.9,beta_2=0.999)`

- Evaluation

- pros
  - computationally efficient and has little memory requirement
  - recommended as a default optimization algorithm
- cons
  - it focus on fast computation speed, in some situations, Adam leads to unconvergence, while some algorithms(SGD) focus on data point.
  - the low learning rate at final may lead to missing the optimum solution.

- AdaGrad**(Adaptive Gradient)



- use different learning rates for parameters

- The more parameters get change, the more minor the learning rate changes. For sparse features, lr needed to be higher than that of dense features.

- Because the occurrence frequency of sparse features is lower.

- Function of weights update

- $w_t = w_{t-1} - \eta'_t \frac{\delta L}{\delta w_{t-1}}$

- where  $\eta'_t = \frac{\eta}{\sqrt[2]{\alpha_t + \epsilon}}$

- $\alpha_t$ : learning rate at time t

- $\eta$ : initial learning rate

- Evaluation

- downside

- decrease  $\alpha$  aggressively and monotonically, the model will stop learning too early, cuz the learning rate is almost close to 0, the accuracy of the model is getting down.

- The initial learning rate should be decided manually

- Pros

- It is fitted for the reality that a dataset contains both sparse features and dense features

## • AdaDelta

- Without sum up the past squared gradients as AdaGrad, we restrict the window size, with exponential weighted average

- Function: the same as AdaGrad, except replace  $\alpha$  with exponential weighted average of squared gradients.

- $\eta'_t = \frac{\eta}{\sqrt[2]{S_{dw_t} + \epsilon}}$

- $S_{dw_t} = \beta S_{dw_{t-1}} + (1 - \beta) \left( \frac{\delta L}{\delta w_{t-1}} \right)^2$

- typically  $\beta = 0.9$  or  $0.95$

- Evaluation

- Pros

- solve the radically diminishing learning rates.

## • Root Mean Square Propagation(RMSProp)

- RPPROP

- solve the problem that some gradients are small while some are huge

- use the sign of the gradient adapting the step size individually for each weight

- Select two gradient and compare their sign


- sign1=sign2→go in right direction, increase the step size

- sign1 ≠ sign2→decrease the step size

- Problem

- does not work well for large dataset and mini-batch updates--> the purpose of RMSProp
- Choose different learning rate for parameters
- accelerating the optimization process by decreasing the # of function evaluations to reach local minimum
- accumulating only the gradients from the most recent iterations
- Functions (for weight  $w$  and bias  $b$ )
  - $v_t = \beta v_{t-1} + (1 - \beta) g_t^2$
  - $w_{t+1} = w_t - \frac{\eta}{\sqrt{v_t + \epsilon}} * g_t$
- Evaluations
  - Pros
    - reduce the monotonically decreasing learning rate
    - quickly converge
    - require less tuning
  - Cons
    - the learning rates should be defined manually
- Coding: `optimizer=keras.optimizers.RMSprop(lr=0.001, rho=0.9)`

## • NAG(Nesterov Accelerated Gradient)



(a) Momentum-Based Gradient Descent

(b) Nesterov Accelerated Gradient Descent

$$\text{Green circle} \Rightarrow \frac{\partial L}{\partial w_0} = \frac{\text{Negative}(-)}{\text{Positive}(+)} \quad \text{Red circle} \Rightarrow \frac{\partial L}{\partial w_0} = \frac{\text{Negative}(-)}{\text{Negative}(-)}$$

- Compared with momentum optimization method, the direction is different, it not at the local position, but at the direction of momentum direction(toward the optimum)
- Function
  - $m \leftarrow \beta m - \eta \delta_\theta L(\theta + \beta m)$  ;the gradient is optimum at  $\theta + \beta m$  not at  $\theta$
  - $\theta = \theta + m$
- coding: `optimizer=keras.optimizer.SGD(lr=0.001,momentum=0.9,nesterov=True)`
- Evulation
  - NAG ends up slight faster than momentum method
  - help reduce oscillations

## • Nadam

- Converge slightly faster than Adam
- Nadam= Adam+Nesterov
- References:
  - A comprehensive guide on deep learning optimizers.<https://www.analyticsvidhya.com/blog/2021/10/a-comprehensive-guide-on-deep-learning-optimizers/>

以上内容整理于 [幕布文档](#)